

# COIT20277 Introduction to Artificial Intelligence

## Week 2 - Lecture

- Machine Learning Overview
- Supervised Learning: Classification



BE WHAT **YOU** WANT TO BE  
**cqu.edu.au**

# Acknowledgement of Country

I respectfully acknowledge the Traditional Custodians of the land on which we live, work and learn. I pay my respects to the First Nations people and their Elders, past, present and future



BE WHAT **YOU** WANT TO BE  
**cqu.edu.au**

# Acknowledgment

The content of this lecture has been adopted from the following book:

- Artificial Intelligence Programming with Python - From Zero to Hero, 2022, Perry Xiao, *John Wiley & Sons, Inc.*, ISBN 978-1-119-82086-4.
- Chapter 3 (Sections 3.1 and 3.2)



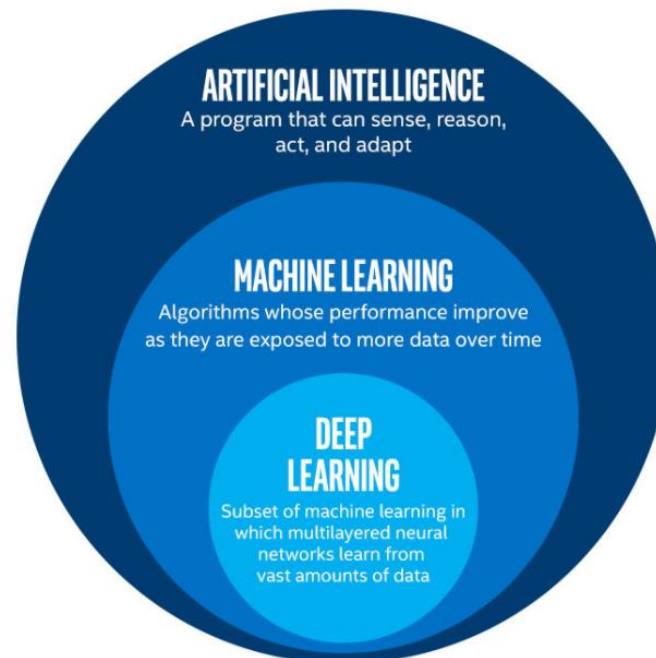
BE WHAT **YOU** WANT TO BE  
**cqu.edu.au**

# Outline

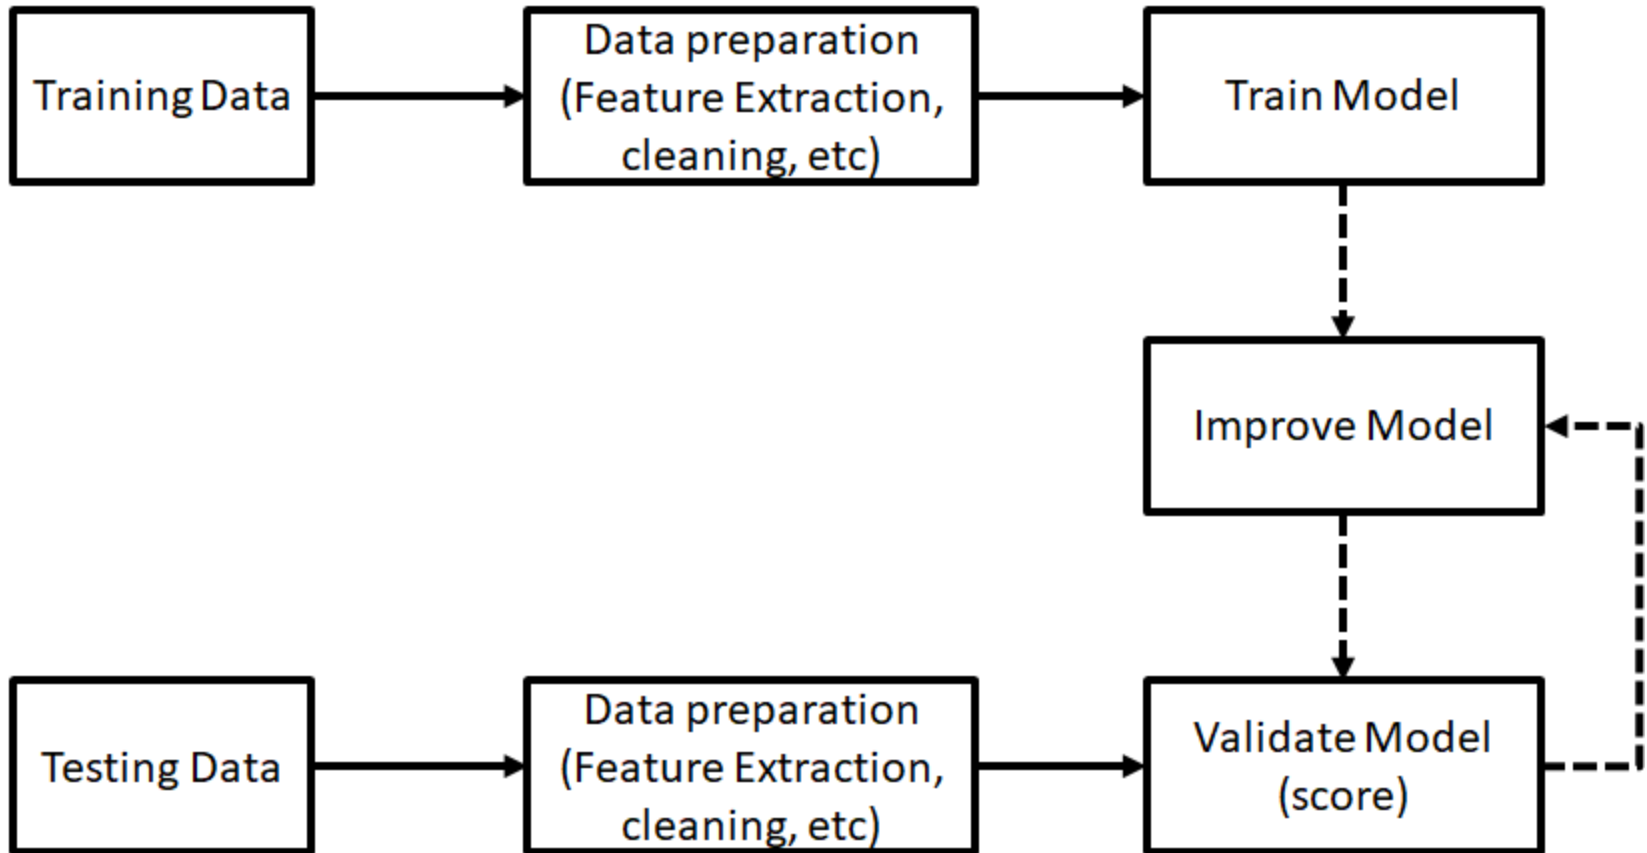
- What is Machine Learning?
- History of Machine Learning
- Types of Machine Learning
- Applications of Machine Learning
- Supervised Learning: Classifications
- Popular supervised learning algorithms
- Ready-made datasets in Scikit-Learn
- Support Vector Machines (SVM)
- Naïve Bayes
- Decision Trees and Random Forests
- K-Nearest Neighbors (K-NN)

# What is Machine Learning (ML)?

- Machine learning is a subset of AI.
- It involves teaching computers to learn from data and analyze data automatically, without human intervention.
- It includes a set of mathematical algorithms that can make decisions or predict results for a given set of data.

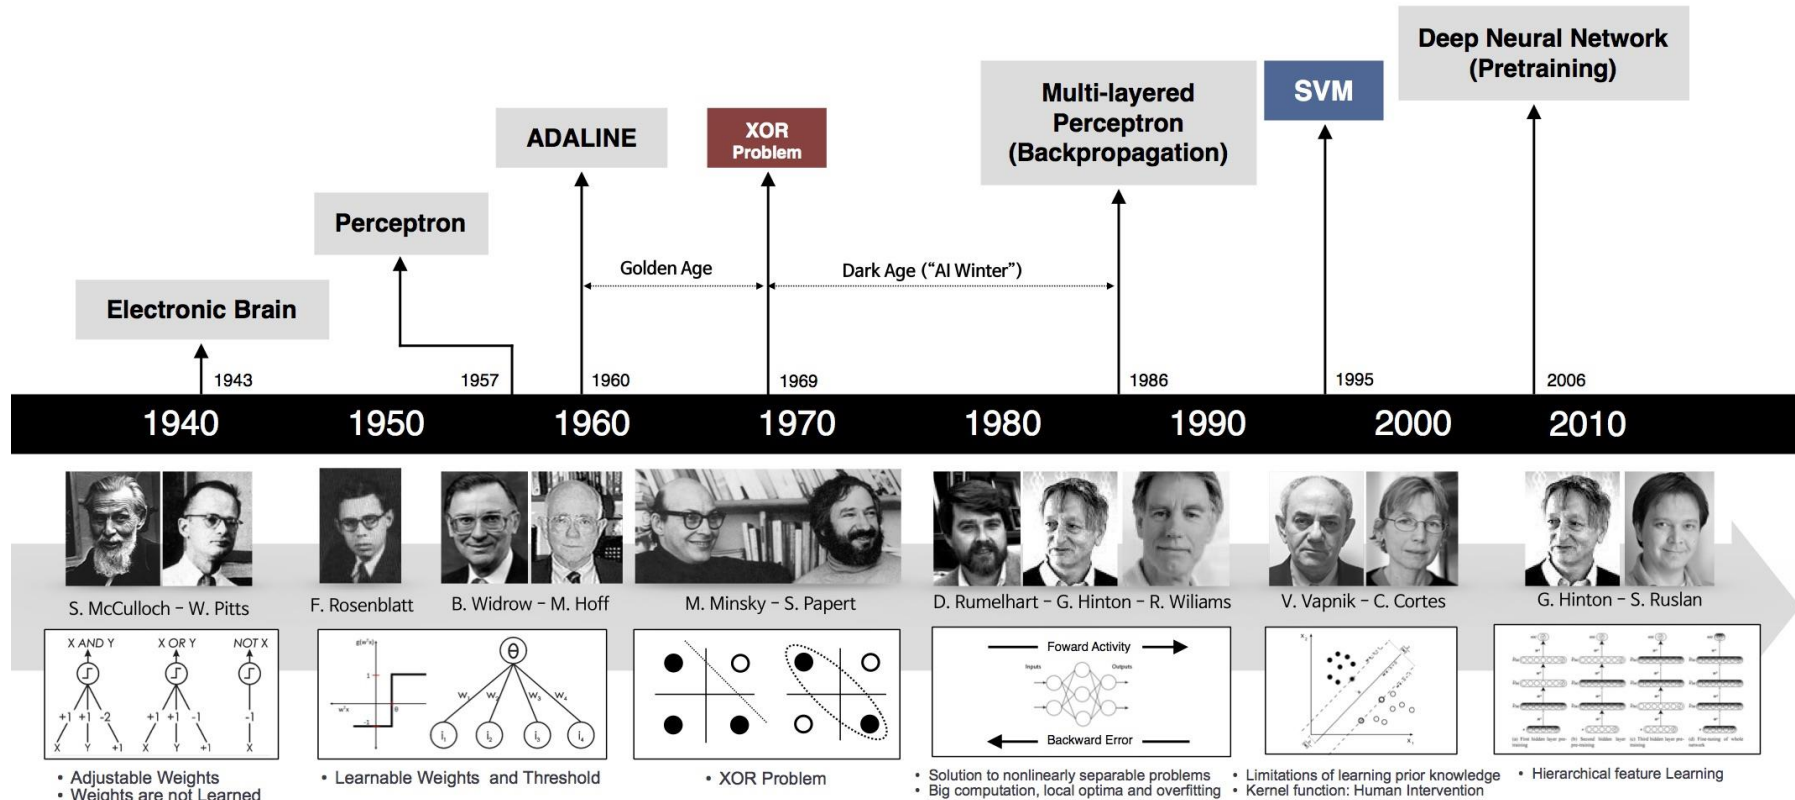


# Machine Learning Pipelines



[This Photo](#) by Unknown Author is licensed under [CC BY-SA](#)

# History of Machine Learning



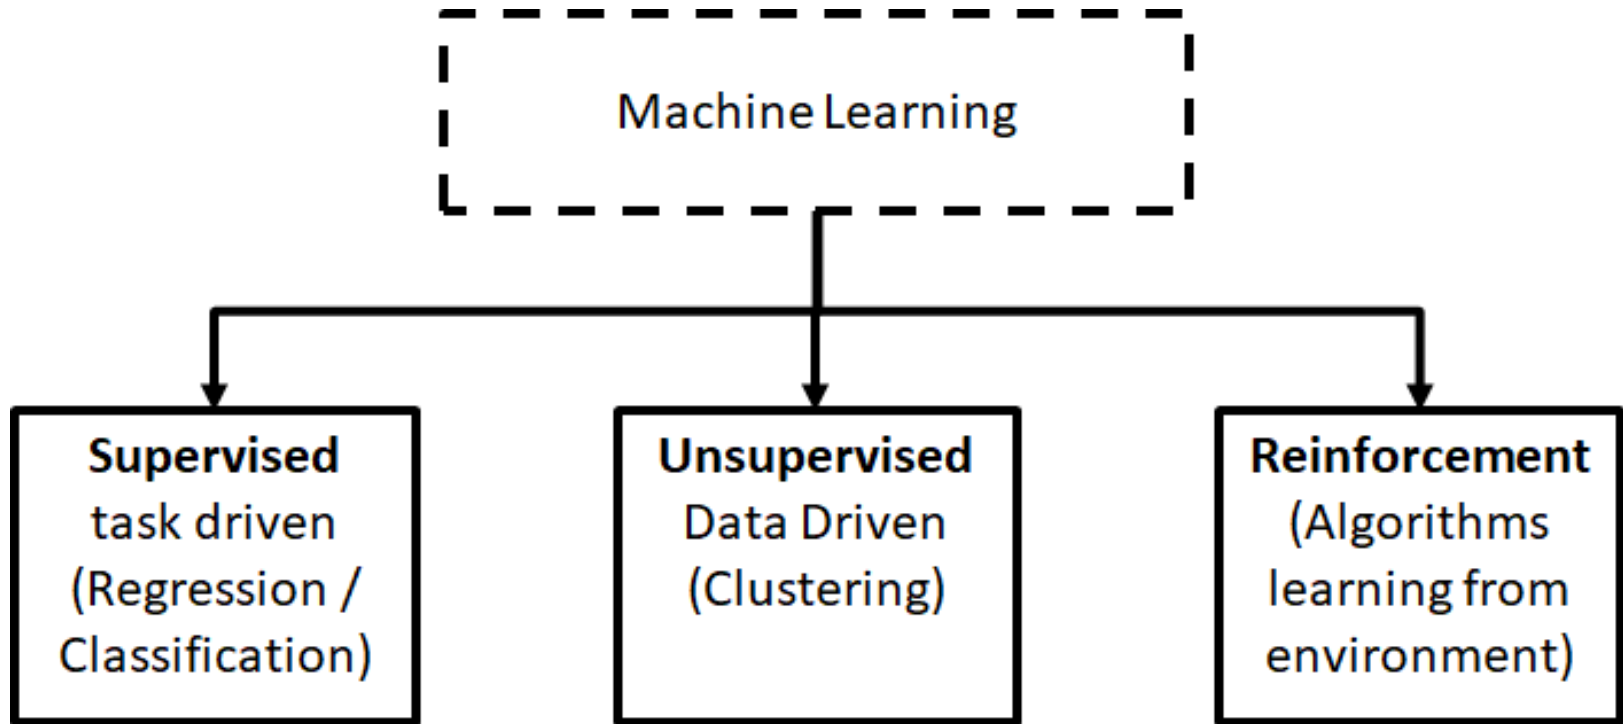
This Photo by Unknown Author is licensed under [CC BY](https://creativecommons.org/licenses/by/4.0/)

# Current Landscape of Machine Learning

- ***Widely adopted across industries:*** Machine learning is being used in diverse sectors like healthcare, finance, and retail to enhance decision-making and customer experience.
- ***Dominance of deep learning:*** Techniques such as CNNs and RNNs are prevalent, especially in image recognition, natural language processing, and speech recognition.
- ***Ethical considerations:*** Growing awareness of bias, fairness, privacy, and accountability is influencing the development and deployment of AI systems.
- ***Interdisciplinary collaboration:*** Collaboration between experts from various fields is common, promoting holistic approaches to AI development and deployment.
- ***Democratization of AI:*** Access to machine learning tools and platforms is becoming more widespread, enabling individuals and organizations to build AI solutions with less expertise.

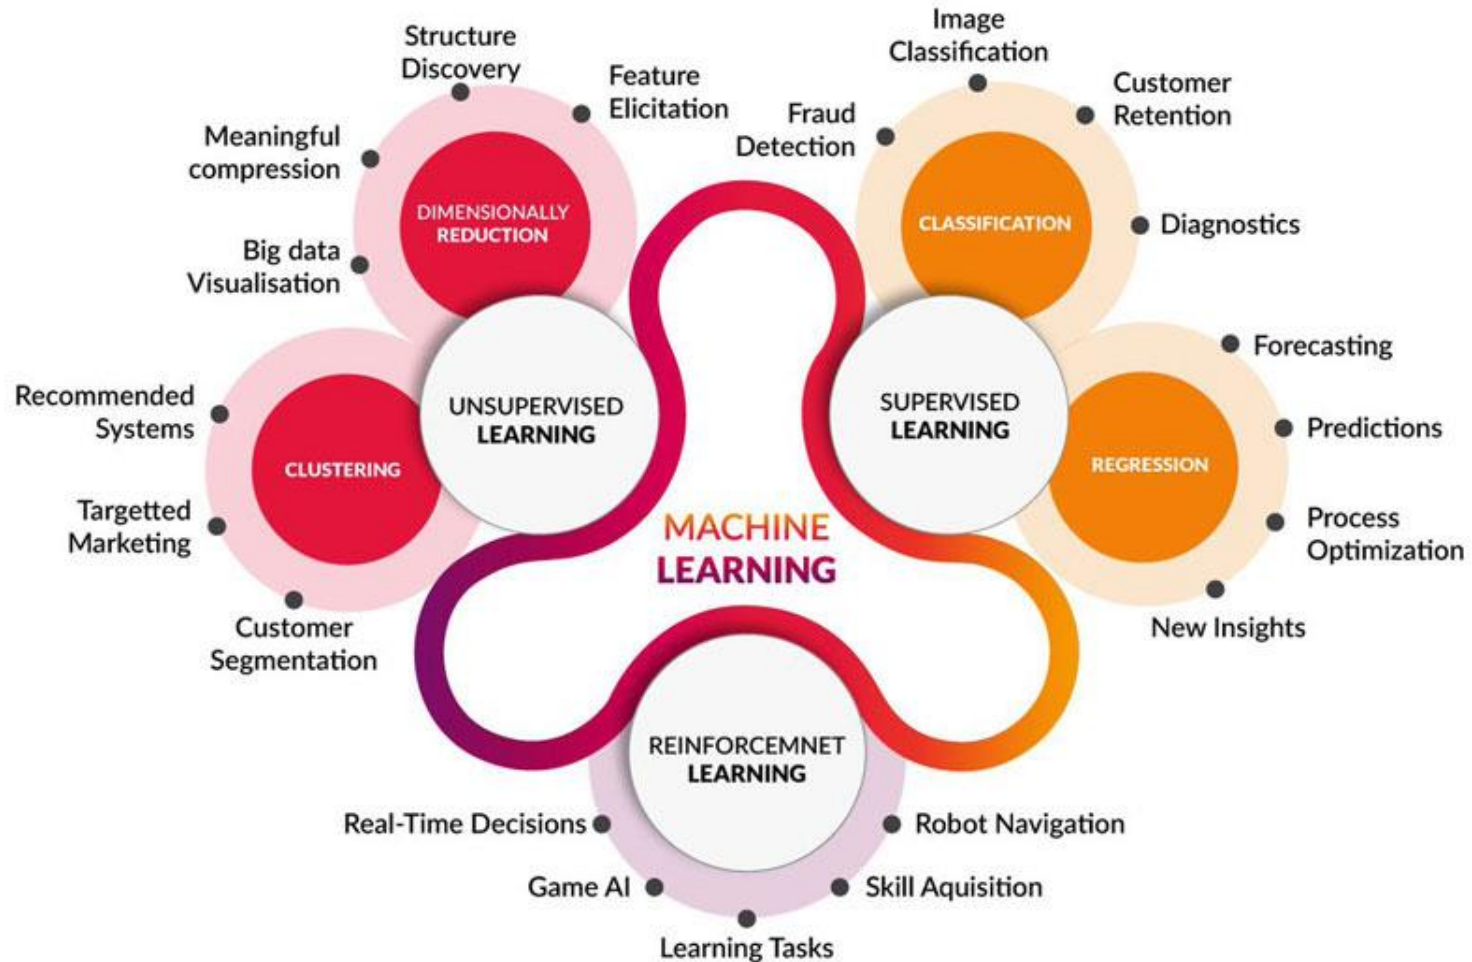


# Types of Machine Learning



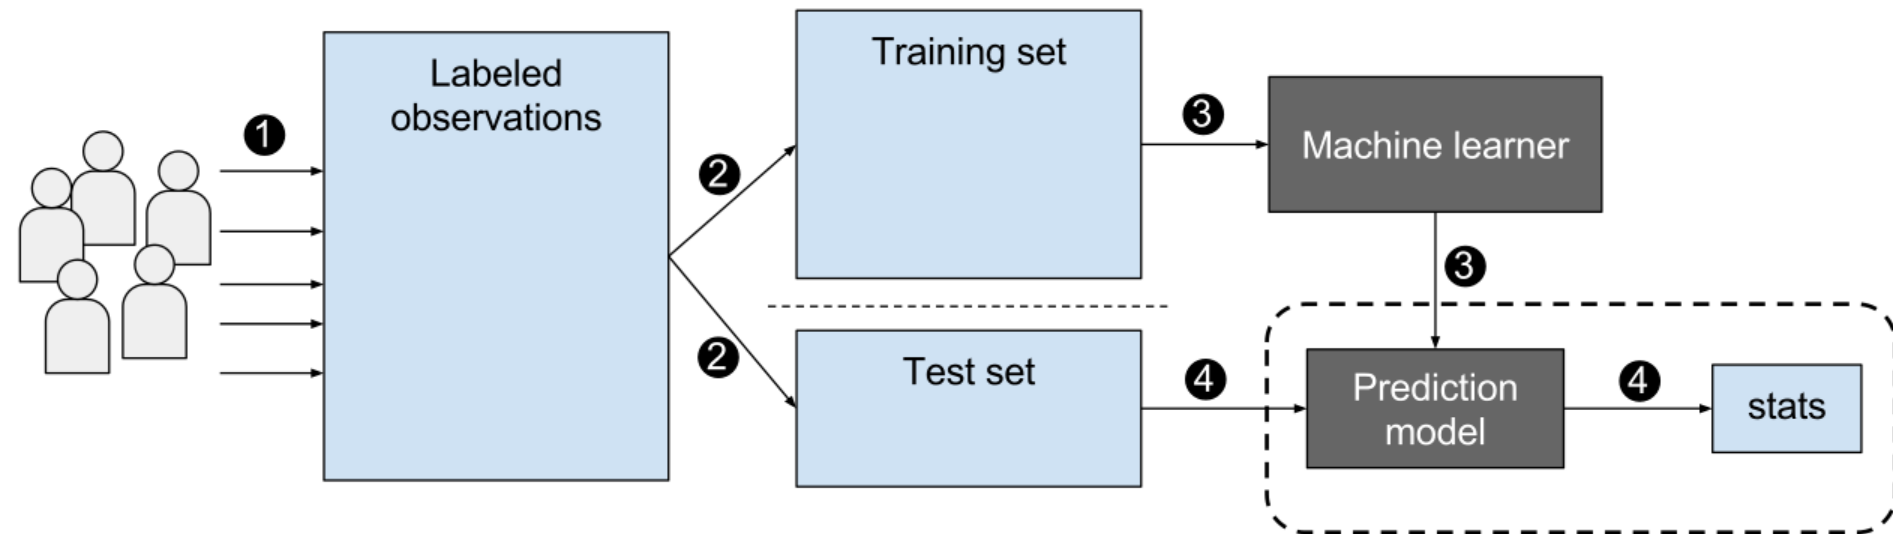
[This Photo](#) by Unknown Author is licensed under [CC BY-SA](#)

# Applications of Machine Learning




# Supervised Learning

- Supervised learning is the most important part of machine learning, used for both classification and regression.
- Classification focuses on predicting the category a sample belongs to.
- Key terms: classes (categories), features (measurements), samples (data points), and parameters (model variables).



# Classification Example


- **Spam email filtering** utilizes supervised learning to categorize emails into spam or non-spam categories.
- Data is labeled as either spam or non-spam, serving as the training set for the classification model.



[This Photo](#) by Unknown Author is licensed under [CC BY-SA](#)

# Classification Example (cont...)

- **EEG signals classification** involves categorizing brain wave patterns recorded through electroencephalography (EEG) into specific classes.
- Using supervised learning, EEG signals are associated with corresponding outputs (class labels).



# Popular Algorithms

- Support Vector Machines (SVM)
- Naïve Bayes
- Linear Discriminant Analysis (LDA)
- Principal Component Analysis (PCA)
- Decision Trees
- Random Forest
- K-Nearest Neighbors (K-NN)
- Artificial Neural Networks (Multilayer Perceptron)
- Find more details and examples using Scikit-Learn:  
[https://scikit-learn.org/stable/supervised\\_learning.html](https://scikit-learn.org/stable/supervised_learning.html)

# What is Scikit-learn?



- Scikit-learn: Open-source Python library for machine learning.
- Features: Comprehensive, user-friendly, efficient, integrates with other Python libraries.
- Algorithms: Supervised and unsupervised learning algorithms included.
- Applications: Data preprocessing, model evaluation, real-world tasks like predictive modeling.
- Community: Active, with extensive documentation and tutorials available.

[This Photo](#) by Unknown  
Author is licensed under  
[CC BY-SA](#)

# Scikit-Learn Datasets




[This Photo](#) by Unknown  
Author is licensed under  
[CC BY-SA](#)

- Ready-made datasets are available for easy access:
  - Toy datasets (e.g., iris flowers, breast cancer).
  - Real-world datasets (e.g., forest cover types).
  - Generated datasets.
  - Find more details: <https://scikit-learn.org/stable/datasets.html>
- Things to keep in mind:
  - Toy datasets may not capture all complexities of real-world problems.
  - Real-world datasets may require further preprocessing and exploration.
  - Always consider the specific research question and data availability when choosing datasets.




# Support Vector Machine (SVM)

- SVM can work on both classification and regression problems.
- Proposed by Vapnik at AT&T Bell Labs in 1963, it uses a statistical learning framework.
- Two-category classification with SVM, separating data points using a hyperplane (straight line in 2-D).
- SVM adjusts the hyperplane to maximize the margin between data points.



# Linear SVM

- All hyperplanes in  $\mathbb{R}^d$  are parameterised by a vector ( $\mathbf{w}$ ) and a constant  $b$ .
- Can be expressed as  $\mathbf{w} \cdot \mathbf{x} + b = 0$  (remember the equation for a hyperplane from algebra!)
- Our aim is to find such a hyperplane  $f(\mathbf{x}) = \text{sign}(\mathbf{w} \cdot \mathbf{x} + b)$ , that correctly classify our data.




Data:  $\langle \mathbf{x}_i, y_i \rangle$ ,  $i=1, \dots, l$

$\mathbf{x}_i \in \mathbb{R}^d$

$y_i \in \{-1, +1\}$


# Linear SVM (Definition)

- Define the hyperplane  $H$  such that:
  - $\mathbf{w} \cdot \mathbf{x}_i + b \geq +1$  when  $y_i = +1$
  - $\mathbf{w} \cdot \mathbf{x}_i + b \leq -1$  when  $y_i = -1$
- $H1$  and  $H2$  are the planes:
  - $H1: \mathbf{w} \cdot \mathbf{x}_i + b = +1$
  - $H2: \mathbf{w} \cdot \mathbf{x}_i + b = -1$
- The points on the planes  $H1$  and  $H2$  are the **Support Vectors**.
- $d^+$  = the shortest distance to the closest positive point, while  $d^-$  = the shortest distance to the closest negative point.
- The **margin** of a separating hyperplane is  $d^+ + d^-$



# Linear SVM (Maximise the Margin)

- We want a classifier with as big a margin as possible.
- Recall the distance from a point  $(x_0, y_0)$  to a line of the form  $Ax + By + c = 0$  is:  
$$|Ax_0 + By_0 + c| / \sqrt{A^2 + B^2}$$
- The distance between H and H1 is:  $|\mathbf{w} \cdot \mathbf{x} + b| / \text{norm}(\mathbf{w}) = 1 / \text{norm}(\mathbf{w})$
- The distance between H1 and H2 is therefore  $2 / \text{norm}(\mathbf{w})$
- In order to maximise the margin, we need to minimize  $\text{norm}(\mathbf{w})$ .



# Solve by a Constrained Optimisation Problem

- Minimise  $\|\mathbf{w}\| = \langle \mathbf{w} \cdot \mathbf{w} \rangle$  subject to  $y_i(\langle \mathbf{x}_i \cdot \mathbf{w} \rangle + b) \geq 1$  for all  $i$ .
- Lagrangian method:  $\max \inf_{\mathbf{w}} L(\mathbf{w}, b, \alpha)$  where  $L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_i \alpha_i [(y_i(\mathbf{x}_i \cdot \mathbf{w}) + b) - 1]$
- At the extremum, the partial derivative of  $L$  with respect to both  $\mathbf{w}$  and  $b$  must be 0.
- Taking the derivatives, and setting them to 0, then substituting back into  $L$  and simplifying, yields:  
$$\text{Maximise } \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j \langle \mathbf{x}_i \cdot \mathbf{x}_j \rangle \text{ subject to } \sum_i y_i \alpha_i = 0 \text{ and } \alpha_i \geq 0$$
- This is an instance of a positive, semi-definite programming problem which can be solved in  $O(n \times \log n)$  time.

# Naïve Bayes

- In 1763, Reverend Thomas Bayes, a mathematician and Presbyterian minister, posthumously published "An Essay towards solving a Problem in the Doctrine of Chances".
- This essay introduced the foundation of what we now call **Bayes Theorem**, offering “a framework for updating beliefs based on new evidence”.




Image courtesy of LANL

[This Photo](#) by Unknown Author is licensed under [CC BY-NC-ND](#)

# The Bayes Classifier

- Problem statement:
  - Given features  $X_1, X_2, \dots, X_n$
  - Predict a label  $Y$
- A good strategy is to predict:

$$\arg \max_Y P(Y | X_1, X_2, \dots, X_n)$$

- For example: *What is the probability that the image represents a dog given its pixels?*
- How do we compute that?

# The Bayes Classifier (cont...)

- Use Bayes Theorem:

Posterior

Likelihood

Prior

$$P(Y|X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n|Y)P(Y)}{P(X_1, \dots, X_n)}$$

Normalisation  
Constant

- To classify if an image is a dog, we first compute the following two probabilities:

$$\begin{aligned} P(Y = \text{Dog}|X_1, \dots, X_n) \\ &= \frac{P(X_1, \dots, X_n|Y = \text{Dog})P(Y = \text{Dog})}{P(X_1, \dots, X_n|Y = \text{Dog})P(Y = \text{Dog}) + P(X_1, \dots, X_n|Y = \neg \text{Dog})P(Y = \neg \text{Dog})} \end{aligned}$$

$$\begin{aligned} P(Y = \neg \text{Dog}|X_1, \dots, X_n) \\ &= \frac{P(X_1, \dots, X_n|Y = \neg \text{Dog})P(Y = \neg \text{Dog})}{P(X_1, \dots, X_n|Y = \text{Dog})P(Y = \text{Dog}) + P(X_1, \dots, X_n|Y = \neg \text{Dog})P(Y = \neg \text{Dog})} \end{aligned}$$

- Classify the image is a dog if  $P(Y = \text{Dog}|X_1, \dots, X_n) \geq P(Y = \neg \text{Dog}|X_1, \dots, X_n)$



# The Bayes Classifier (cont...)

- For the Bayes Classifier, we need to learn two functions, the *likelihood* and the *prior*.
- The problem with explicitly modelling  $P(X_1, \dots, X_n|Y)$  is that there are usually way too many parameters.
- We'll run out of space.
- We'll run out of time.
- And we'll need a lot of training data (which is usually not available).
- The solution lies in the Naïve Bayes Assumption.

# The Naïve Bayes Model

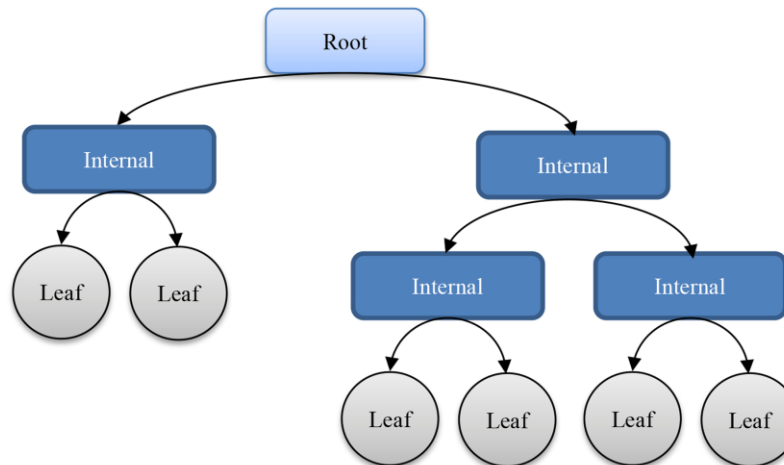
- The Naïve Bayes Assumption: Assume that all features  $X_1, X_2, \dots, X_n$  are ***independent*** given the class label  $Y$ .
- With this assumption, the following equation holds:

$$P(X_1, \dots, X_n | Y) = \prod_{i=1}^n P(X_i | Y)$$

- Without the ***independent*** assumption, the number of parameters for modelling  $P(X_1, \dots, X_n | Y)$  is  $2(2^n - 1)$ .
- With the ***independent*** assumption, the number of parameters is reduced to  $2n$ .

# Decision Trees

- Tree-like structures where each node represents a feature (question) and branches represent possible answers.
- Each branch leads to a new node, asking another question or reaching a leaf node containing the final prediction.
- They are easy to understand and visualise.
- Used for both classification (predicting categories) and regression (predicting continuous values).



[This Photo](#) by Unknown Author is licensed under [CC BY](#)

# Decision Trees (cont...)

- Decision trees are nonparametric, supervised learning methods.
- Represented as an upside-down tree, where derived rules are used to guide decisions.
- Nodes represent query variables, edges represent their values, branches represent if-then rules.
- Deeper trees yield more complex rules and models.
- **Goal:** Training algorithm creates decision tree based on dataset, produces rules for prediction.

# Strengths of Decision Trees

- ***Interpretability***: We can easily understand the thought process behind each prediction due to the tree structure.
- ***Ability to handle mixed data types***: Decision trees can handle both numerical and categorical features without complex preprocessing.
- ***Robustness to missing data***: They can impute missing values by following the most common branch for that feature.

# Weaknesses of Decision Trees

- ***Overfitting***: Decision trees can become too specific to the training data, leading to poor performance on unseen data.
- ***Instability***: Small changes in the data can lead to significant changes in the tree structure and predictions.
- ***High variance***: Individual trees can be sensitive to changes in the training data, leading to inconsistent predictions.

# Random Forest

- Random forests are ensembles of decision trees, combining multiple trees for improved accuracy and stability.
- Overfitting: Random forest reduces overfitting and improves performance by combining output of individual decision trees.
- Each tree in the forest is trained on a different bootstrap sample of the data and uses a subset of features randomly selected without replacement.
- The final prediction is based on the majority vote (for classification) or the average (for regression) of the individual tree predictions.

# Strengths of Random Forest

- ***Improved accuracy***: Random forests often outperform individual decision trees due to reduced overfitting and variance.
- ***Robustness to noise and outliers***: By averaging predictions from multiple trees, random forests are less sensitive to outliers and noise in the data.
- ***Ability to handle high-dimensional data***: They can effectively deal with datasets with many features without significant performance degradation.

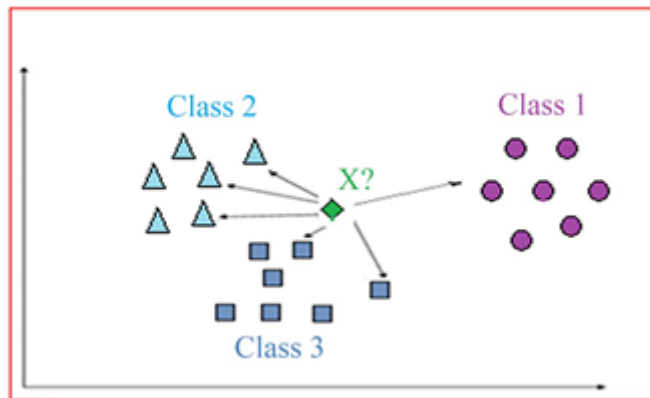


# K-Nearest Neighbors (K-NN)

- K-NN Algorithm: Utilises K nearest points to determine classification or regression.
- K-nearest neighbors are different from K-means, which is for clustering.
- K-NN classifies data points based on their "neighborhood" in the feature space.
- Imagine data points as colored dots on a map. Their colors represent their class (e.g., red for apples, green for oranges).
- When encountering a new data point, K-NN finds its K nearest neighbors (say, 3 closest dots).
- The new point's class is assigned based on the majority vote of its neighbors (e.g., if 2 neighbors are red, the new one is likely red too).

# Choosing the Right K: A Balancing Act

- K, the number of neighbors, significantly impacts K-NN's performance.
- K too low (e.g.,  $K=1$ ) can be sensitive to noise and outliers, leading to overfitting.
- K too high (e.g.,  $K=20$ ) can smooth out details, causing underfitting.
- Finding the optimal K often involves trial and error or techniques like cross-validation.



[This Photo](#) by Unknown Author is licensed under [CC BY](#)

# Strengths and Applications of K-NN

- Strengths:
  - Simple to understand and implement.
  - Effective for small datasets and high-dimensional data.
  - Can handle mixed data types (numerical and categorical).
- Applications:
  - Image classification (recognizing handwritten digits, categorizing photos).
  - Customer segmentation (grouping customers based on purchase history).
  - Fraud detection (identifying unusual transactions).
  - Spam filtering (classifying emails as spam or not spam).



# THANK YOU

TIME FOR DISCUSSION & QUESTIONS