

COIT20277 Introduction to Artificial Intelligence

Week 3

- Supervised Learning: Regression
- Unsupervised Learning



BE WHAT **YOU** WANT TO BE
cqu.edu.au

Acknowledgement of Country

I respectfully acknowledge the Traditional Custodians of the land on which we live, work and learn. I pay my respects to the First Nations people and their Elders, past, present and future



BE WHAT **YOU** WANT TO BE
cqu.edu.au

Supervised Learning: Regressions

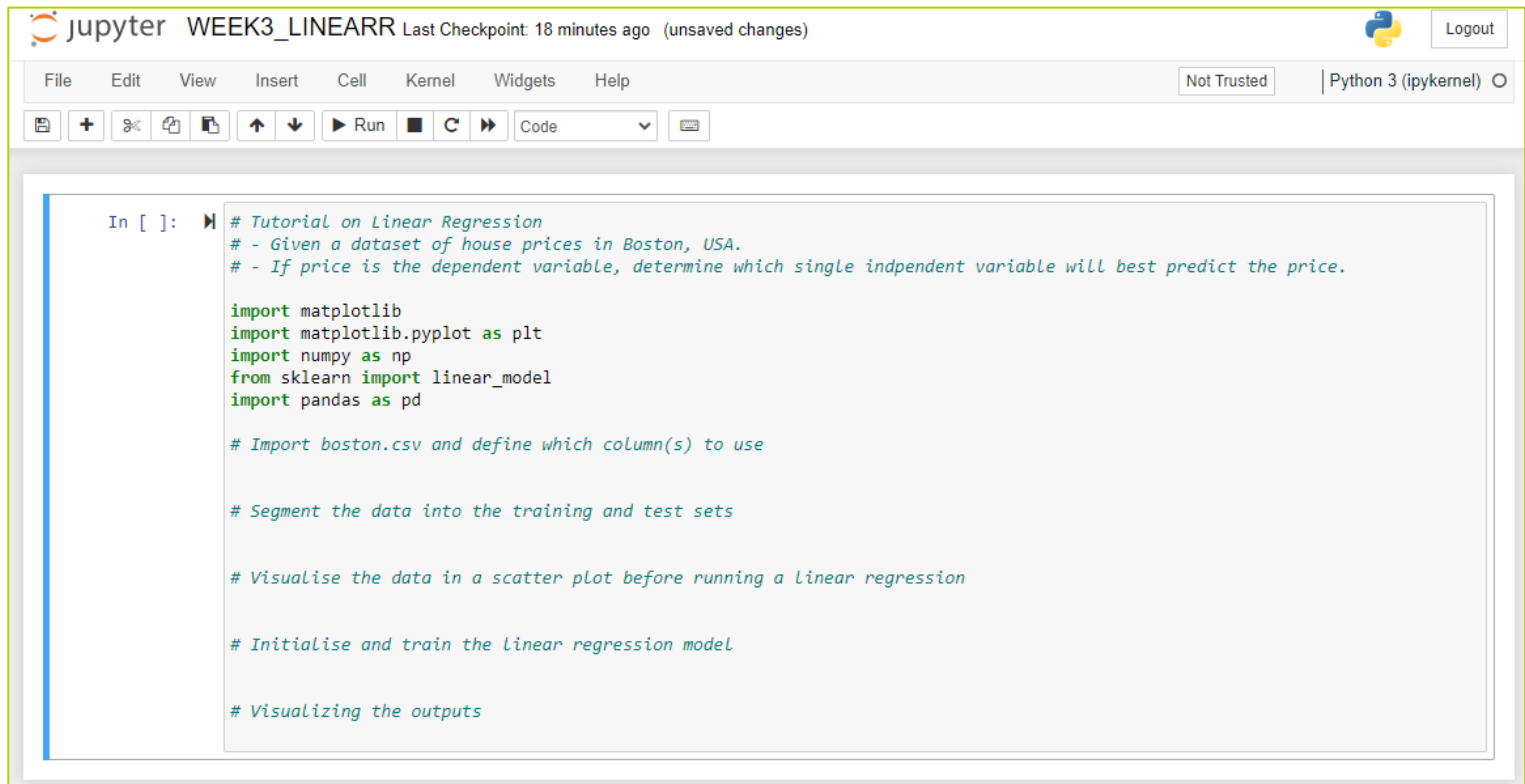
- The supervised learning algorithms explored in this tutorial are:
 - Linear Regression
 - Logistic Regression
- You can learn more about these algorithms implemented in scikit-learn by accessing https://scikit-learn.org/stable/supervised_learning.html and https://scikit-learn.org/stable/unsupervised_learning.html

Unsupervised Learning

- The unsupervised learning algorithm explored in this tutorial is:
 - K-means clustering
- You can learn more about K-means and other unsupervised learning algorithms implemented in scikit-learn by accessing https://scikit-learn.org/stable/unsupervised_learning.html

Linear Regression

- Tutorial on Linear Regression
 - Given a dataset of house prices, `boston.csv`, in Boston, USA.
 - Price is the dependent variable, while other columns are independent variables.
 - Your job is to complete the Python program, `WEEK3_LINEAR.ipynb` to determine which independent variable will best predict the price.



The screenshot shows a Jupyter Notebook titled "WEEK3_LINEAR" with a "Last Checkpoint: 18 minutes ago (unsaved changes)" status. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help), a "Not Trusted" security warning, and a "Python 3 (ipykernel)" environment indicator. The code cell contains the following Python code:

```
In [ ]: # Tutorial on Linear Regression
# - Given a dataset of house prices in Boston, USA.
# - If price is the dependent variable, determine which single independent variable will best predict the price.

import matplotlib
import matplotlib.pyplot as plt
import numpy as np
from sklearn import linear_model
import pandas as pd

# Import boston.csv and define which column(s) to use

# Segment the data into the training and test sets

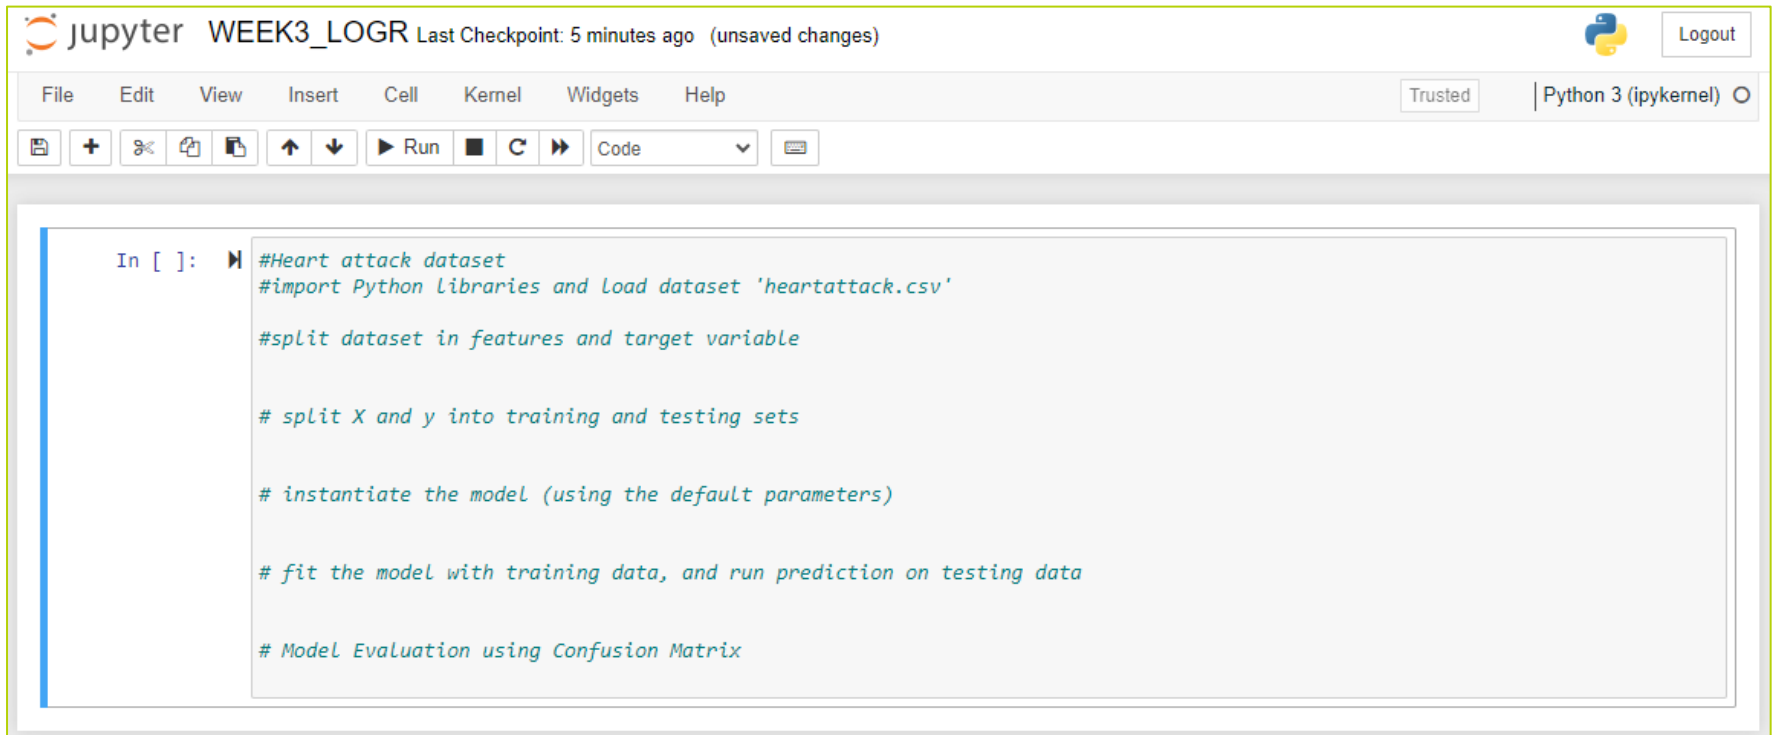
# Visualise the data in a scatter plot before running a linear regression

# Initialise and train the linear regression model

# Visualizing the outputs
```

Logistic Regression

- Tutorial on Logistic Regression
 - Given a dataset on persons who had or hadn't had a heart attack before, `heartattack.csv`.
 - The last column is a binary [0,1] meaning a heart attack had occurred or not.
 - Your task is to complete the Python program, `WEEK3_LOGR.ipynb` to train a logistic regressor to predict if a person will have a heart attack or not.

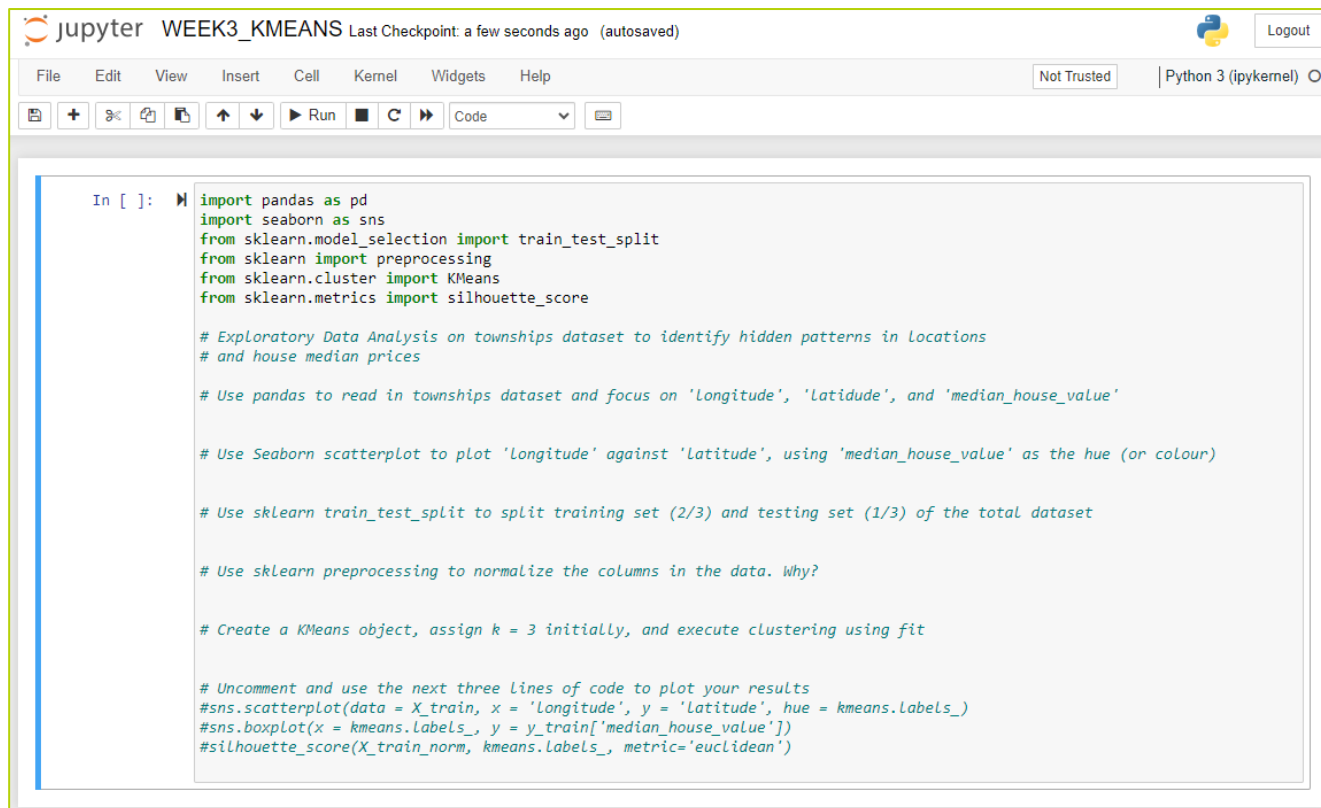


The screenshot shows a Jupyter Notebook titled "WEEK3_LOGR" with a "Last Checkpoint: 5 minutes ago (unsaved changes)" status. The interface includes a top menu bar with options like File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. Below the menu is a toolbar with icons for saving, running, and other actions. The main area displays a Python code cell with the following comments:

```
In [ ]: #Heart attack dataset  
#import Python libraries and load dataset 'heartattack.csv'  
  
#split dataset in features and target variable  
  
# split X and y into training and testing sets  
  
# instantiate the model (using the default parameters)  
  
# fit the model with training data, and run prediction on testing data  
  
# Model Evaluation using Confusion Matrix
```

K-means Clustering

- Tutorial on K-means Clustering
 - Given a dataset on townships in Australia with longitude, latitude, population, median house prices, etc. in `townships.csv`.
 - Your task is to implement K-means Clustering in a Python program similar as below, `WEEK3_KMEANS.ipynb`, to find the hidden clusters.
 - You can experiment with different values for K, the number of clusters.



The screenshot shows a Jupyter Notebook titled "WEEK3_KMEANS" with a "Last Checkpoint: a few seconds ago (autosaved)" status. The interface includes a top bar with "File", "Edit", "View", "Insert", "Cell", "Kernel", "Widgets", and "Help" menus. Below the menus is a toolbar with icons for file operations, a "Run" button, and a "Code" dropdown. The main area contains a code cell with the following Python code:

```
In [ ]: import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

# Exploratory Data Analysis on townships dataset to identify hidden patterns in Locations
# and house median prices

# Use pandas to read in townships dataset and focus on 'Longitude', 'Latitude', and 'median_house_value'

# Use Seaborn scatterplot to plot 'Longitude' against 'Latitude', using 'median_house_value' as the hue (or colour)

# Use sklearn train_test_split to split training set (2/3) and testing set (1/3) of the total dataset

# Use sklearn preprocessing to normalize the columns in the data. Why?

# Create a KMeans object, assign k = 3 initially, and execute clustering using fit

# Uncomment and use the next three lines of code to plot your results
#sns.scatterplot(data = X_train, x = 'Longitude', y = 'Latitude', hue = kmeans.Labels_)
#sns.boxplot(x = kmeans.Labels_, y = y_train['median_house_value'])
#silhouette_score(X_train_norm, kmeans.Labels_, metric='euclidean')
```



THANK YOU

TIME FOR DISCUSSION & QUESTIONS